



# Data Mining for NLP

## 3- Unsupervised Approaches to Explore Data

These slides will be available on Arche

# Course Objective

**Goal:** Use Data Mining methods to improve NLP workflow

- Present methods to **explore textual data**
- Focus on **Machine Learning methods** to deal with NLP data
- Focus on **empirical approaches**



# Course Logistics

- 3 sessions (3 lectures + 3 lab)
- 1h30 lectures
- 3 labs (google colab)

Material on Arche



# Course Evaluation

- **Final Exam**
  - Mix between questions and code completion
  - 30% of the final grade Data Mining
  - Exam 29/01 5pm

# Exploring Data: General Idea

1. **Represent the data:** one-hot, embeddings, features bags, etc.
2. **IF** there are some labels, use them to better understand the data in a **supervised learning** setting
  - a. Feature importance score
  - b. Leave One Out strategy
  - c. Attention scores
3. Use **unsupervised learning** to better understand underlying distribution of the data
  - a. PCA
  - b. Simple unsupervised approaches
  - c. Latent space (auto-encoders)
4. Mix both

# How to represent input tokens?

## Solution 1: 1-Hot Encoding

1. We associate each *token* to a **1-hot vector of size  $D$**

movie =  $[1, 0, \dots, 0, 0, 0]$

hotel =  $[0, 1, \dots, 0, 0, 0]$

...

art =  $[0, 0, \dots, 0, 0, 1]$

1. **Concatenate them** to get a unidimensional vector

# 1-Hot Encoding as inputs

$$dnn_{\theta} : \{0, 1\}^{|V|*K} \rightarrow [0, 1]^V$$
$$x = ([x_1, \dots, x_K]) \mapsto \hat{p}$$

- First hidden layer is of size  $|V|*K$
- Taking as input **a sparse vector**

# 1-Hot Encoding as inputs

$$\begin{aligned} dnn_{\theta} : \quad & \{0, 1\}^{|V|*K} \rightarrow [0, 1]^V \\ & x = ([x_1, \dots, x_K]) \mapsto \hat{p} \end{aligned}$$

**First hidden layer:**

*assuming tanh as the activation function, dimension  $\delta$*

$$h_1 = \tanh(W.x) \text{ s.t. } W \in \mathbb{R}^{\delta \times (|V|*K)}$$



# 1-Hot Encoding as inputs

$$dnn_{\theta} : \{0, 1\}^{|V|*K} \rightarrow [0, 1]^V$$
$$x = ([x_1, \dots, x_K]) \mapsto \hat{p}$$

## Limits

- The representation of each token is fixed and a 1-hot vector
- In this approach, **we do not learn** a representation of **each input token**

# How to represent input tokens?

**Solution 2: Integrate a Dense Embedding Layer**

# How to represent input tokens?

## Solution 2: Integrate a Dense Embedding Layer

We define a dense embedding layer  $E \in \mathbb{R}^{\delta_e \times |V|}$ .

This means that for each token  $t \in V$  indexed by  $j$  in the vocabulary  $V = \{t_1, \dots, t_{|V|}\}$  we have  $t_j$  *embedded* by the vector  $E_{\cdot j}$  (i.e. column of the matrix  $E$  indexed by  $j$ ) of dimension  $\delta_e$  (the dimension of the embedding vectors).

# How to represent input tokens?

## Solution 2: Integrate a Dense Embedding Layer

We define a dense embedding layer  $E \in \mathbb{R}^{\delta_e \times |V|}$ .

This means that for each token  $t \in V$  indexed by  $j$  in the vocabulary  $V = \{t_1, \dots, t_{|V|}\}$  we have  $t_j$  embedded by the vector  $E_{.j}$  (i.e. column of the matrix  $E$  indexed by  $j$ ) of dimension  $\delta_e$  (the dimension of the embedding vectors).

- **$E$  is part of the parametrization** of the model like any other layers
- We can train it during **backprop end-to-end**

# Dense Embedding Layer

$$dnn_{\theta} : \mathbb{R}^{|K| * \delta_e} \rightarrow [0, 1]^V$$

$$x = ([x_1, \dots, x_K]) \mapsto \hat{p}$$

s.t.  $x_i = E_{.j} \in \mathbb{R}^{\delta_e}$  with token  $t_i$  indexed by  $j$  in  $V$

# Dense Embedding Layer

$$dnn_{\theta} : \mathbb{R}^{|K| * \delta_e} \rightarrow [0, 1]^V$$

$$x = ([x_1, \dots, x_K]) \mapsto \hat{p}$$

s.t.  $x_i = E_{.j} \in \mathbb{R}^{\delta_e}$  with token  $t_i$  indexed by  $j$  in  $V$

# Dense Embedding Layer

$$dnn_{\theta} : \mathbb{R}^{|K| * \delta_e} \rightarrow [0, 1]^V$$

$$x = ([x_1, \dots, x_K]) \mapsto \hat{p}$$

s.t.  $x_i = E_{.j} \in \mathbb{R}^{\delta_e}$  with token  $t_i$  indexed by  $j$  in  $V$

- $E$  is a dense embedding matrix
- We can **learn a representation vector for each token in the vocabulary**

# Why is an Embedding Layer much better?

**Trainable Dense Embedding layers** are a “game changer” for Deep Learning Models in NLP i.e. Generalization is much better compared to 1-hot encoding

Why?

(intuition)

$t$  and  $t'$  that have the embedding vectors (in  $E$ )  $x$  and  $x'$ . e.g.  $t = \text{"dog"}$  and  $t' = \text{"cat"}$



# Why is an Embedding Layer much better?

**Trainable Dense Embedding layers** are a “game changer” for Deep Learning Models in NLP i.e. Generalization is much better compared to 1-hot encoding

Why?

(intuition)

$t$  and  $t'$  that have the embedding vectors (in  $E$ )  $x$  and  $x'$ . e.g.  $t = \text{"dog"}$  and  $t' = \text{"cat"}$

1. Let's assume that during training token the model has seen much less frequently *cat* than *dog*

# Why is an Embedding Layer much better?

**Trainable Dense Embedding layers** are a “game changer” for Deep Learning Models in NLP i.e. Generalization is much better compared to 1-hot encoding

Why?

(intuition)

$t$  and  $t'$  that have the embedding vectors (in  $E$ )  $x$  and  $x'$ . e.g.  $t = \text{"dog"}$  and  $t' = \text{"cat"}$

1. Let's assume that during training token the model has seen much less frequently **cat** than **dog**
2. But “luckily”  **$x$  and  $x'$  have similar embedding** vectors (i.e  $\cos(x, x') \sim 1$ )

# Why is an Embedding Layer much better?

**Trainable Dense Embedding layers** are a “game changer” for Deep Learning Models in NLP i.e. Generalization is much better compared to 1-hot encoding

Why?

(intuition)

$t$  and  $t'$  that have the embedding vectors (in  $E$ )  $x$  and  $x'$ . e.g.  $t = \text{"dog"}$  and  $t' = \text{"cat"}$

1. Let's assume that during training token the model has seen much less frequently **cat** than **dog**
2. But “luckily”  **$x$  and  $x'$  have similar embedding** vectors (i.e  $\cos(x, x') \sim 1$ )
3. **When the model *dnn* sees**, at test time, **cat** it will be likely to model **dog** much better than in a 1-hot modeling case **by using this similarity**

# How to initialize an Embedding Layer?

Similarly to all other parameters in a deep learning model

- Before starting training: **we can simply initialize the embedding matrix randomly**
- Before training, the similarity between **embedding word vectors is random**

# How to initialize an Embedding Layer?

Similarly to all other parameters in a deep learning model

- Before starting training: **we can simply initialize the embedding matrix randomly**
- Before training, the similarity between **embedding word vectors is random**

**Can we do better?**

# How to initialize an Embedding Layer?

Similarly to all other parameters in a deep learning model

- Before starting training: **we can simply initialize the embedding matrix randomly**
- Before training, the similarity between **embedding word vectors is random**

**Can we do better?**

- In previous lecture we have seen how to represent good dense embedding vector with skip-gram word2vec model
- **We can simply initialize our word embedding matrix with word2vec vectors**

# How to initialize an Embedding Layer?

Initializing with a **pretrained embedding** layer was also a *gamechanger* for many NLP tasks and many Deep Learning architecture

## Conditions to use a pretrained embedding layer:

- The token in our vocabulary must be in the training of the word2vec model
- For the one that were not seen, we can simply initialize them randomly
  - For LLMs it is even better to average some parts:
    - <https://nlp.stanford.edu/~johnhew/vocab-expansion.html>

# Transfer Learning in NLP

Initializing with a **pretrained embedding** layer is also a *game changer* for many NLP tasks and many Deep Learning architecture

It is called **Transfer Learning**



# Embedding Layer Summary

- Trainable Dense Embedding Layer are a *game changer* for Deep Learning Models
- Even more when we can use a pretrained embedding layers (e.g. with word2vec)
- They can be used with all Deep Learning Architectures
- **For all NLP tasks**

# General DL Framework

# Framework

We want to model  $(X_1, \dots, X_T)$  i.e. find the correct label  $Y$

$$dnn_{\theta} : \mathbb{R}^{d,T} \rightarrow \mathbb{R}^p \text{ or } [|0, K|]^p$$

$$(X_1, \dots, X_T) \mapsto \hat{Y}$$

- Output space is  $\mathbb{R}^p$  for **Regression** tasks
- Output space is  $[|0, K|]^p$  for **Classification** tasks

# Framework

We want to model  $(X_1, \dots, X_T)$  i.e. find the correct label  $Y$

$$dnn_{\theta} : \mathbb{R}^{d,T} \rightarrow \mathbb{R}^p \text{ or } [|0, K|]^p$$

$$(X_1, \dots, X_T) \mapsto \hat{Y}$$

Questions: **when we do Deep Learning...**

- How do we **define**  $dnn_{\theta}$  ?
- How do we **train**  $dnn_{\theta}$  with data ?

# Framework

Given a sequence of vectors  $(X_1, \dots, X_T)$  we want to predict  $Y$

$$\begin{aligned} dnn_{\theta} : \quad \mathbb{R}^{d,T} &\rightarrow \mathbb{R}^p \text{ or } [0, K]^p \\ (X_1, \dots, X_T) &\mapsto \hat{Y} \end{aligned}$$

Most Deep Learning Models:

- are **parametric** (i.e.  $\theta \in \mathbb{R}^D$  )
- defined as a **composition of “simple” functions (linear & non-linear)**
- are trained in an **end-to-end** fashion with **backpropagation**

NB: In Deep Learning, **the parametrization of  $dnn$**  is called **the Architecture**

# Different Types of Architecture

How can we define our predictive function  $dnn_{\theta}$  ?

- Multi-Layer Perceptron
- Recurrent Layers
- Attention Layers
- Self-Attention Layers (in a Transformer Architecture)

# Different Types of Architecture

How can we define our predictive function  $dnn_{\theta}$  ?

- Multi-Layer Perceptron
- Recurrent Layers
- Attention Layers
- Self-Attention Layers (in a Transformer Architecture)

How do we **train our model**? (i.e. estimate the parameters of the model)

- **Stochastic Gradient Descent** also called **backpropagation** in this context

# Output Activation & Loss

## Softmax Function

$$\text{softmax}(s) = \left( \frac{e^{s_i}}{\sum_k e^{s_k}} \right)_{i \in [1, V]}, \text{ for } s \in \mathbb{R}^V$$

## Loss Function

$$l(p, \hat{p}) = CE(p, \hat{p}) = \sum_{i \in [0, V-1]} p_i \log(\hat{p}_i)$$



# Loss Functions

Based on the task we aim at modeling, we can use:

**For Regression: Mean-Square Error**

$$l(y, \hat{y}) = \|y - \hat{y}\|_2^2 = \sum_i (y_i - \hat{y}_i)^2 \text{ assuming } y_i, \hat{y}_i \in \mathbb{R}$$

**For Classification: Cross-Entropy Loss**

$$l(y, \hat{y}) = CE(y, \hat{y}) = \sum_i y_i \log(\hat{y}_i) \text{ assuming } y_i, \hat{y}_i \in [0, 1]$$

Most NLP tasks will be based on the **Cross-Entropy loss**

# Encoder - Decoder

# Framework

We assume an input sequence of tokens  $(x_1, \dots, x_T) \in V^T$   
a target sequence of tokens  $(y_1, \dots, y_{T'}) \in V'^{T'}$ .

# Framework

We assume an input sequence of tokens  $(x_1, \dots, x_T) \in \underline{V}^T$   
a target sequence of tokens  $(y_1, \dots, y_{T'}) \in \underline{V'}^{T'}$ .

# Framework

We assume an input sequence of tokens  $(x_1, \dots, x_T) \in V^T$   
a target sequence of tokens  $(y_1, \dots, y_{T'}) \in V'^{T'}$ .

**Our goal is to estimate:**

$$p_{\theta}(y_1, \dots, y_{T'} | x_1, \dots, x_T)$$

# Framework

We assume an input sequence of tokens  $(x_1, \dots, x_T) \in V^T$   
a target sequence of tokens  $(y_1, \dots, y_{T'}) \in V'^{T'}$ .

**Our goal is to estimate:**

$$p_{\theta}(y_1, \dots, y_{T'} | x_1, \dots, x_T)$$

**We frame it as a classification task:**

$$\hat{y}_t = \operatorname{argmax}_{y \in V'} p_{\theta}(y | (x_1, \dots, x_T), (y_1, \dots, y_{t-1})) \quad \forall t \in [1, T']$$

# What Architecture?

$$p_{\theta}(y_t | (x_1, \dots, x_T), (y_1, \dots, y_{t-1}))$$

We want to model an output sequence conditioned on an input *sequence*

With deep learning, we do that with: **a encoder-decoder model**

**NB:** also called “sequence to sequence” or “seq2seq”

# The Encoder-Decoder Paradigm

$$p_{\theta}(y_t | (x_1, \dots, x_T), (y_1, \dots, y_{t-1}))$$

## Intuition:

- **We know how to model a single sequence at a time with a DL model**  
*E.g. with a LSTM or a Transformer*
- **Here we want to model two sequences together**  
*One conditioned on the other*



# The Encoder-Decoder Paradigm

$$p_{\theta}(y_t | (x_1, \dots, x_T), (y_1, \dots, y_{t-1}))$$

## Intuition:

- **We know how to model a single sequence at a time with a DL model**  
*E.g. with a LSTM or a Transformer*
  - **Here we want to model two sequences together**  
*One conditioned on the other*
- **Combine two Deep Learning Architectures together**

# The Encoder-Decoder Paradigm

$$p_{\theta}(y_t | (x_1, \dots, x_T), (y_1, \dots, y_{t-1}))$$

**Encode input sequence**

$$\begin{aligned} \text{enc}_{\theta_e} : \quad V^T &\rightarrow \mathbb{R}^T \\ (x_1, \dots, x_T) &\mapsto (h_1, \dots, h_T) \end{aligned}$$

# The Encoder-Decoder Paradigm

$$p_{\theta}(y_t | (x_1, \dots, x_T), (y_1, \dots, y_{t-1}))$$

**Encode input sequence**

$$enc_{\theta_e} : V^T \rightarrow \mathbb{R}^T$$

$$(x_1, \dots, x_T) \mapsto (h_1, \dots, h_T)$$

**Decode target sequence**  
given hidden states of  
the encoder

$$dec_{\theta_d} : \mathbb{R}^T \times V'^t \rightarrow [0, 1]^V$$
$$((h_1, \dots, h_T), (y_1, \dots, y_t)) \mapsto \hat{p}$$

# The Encoder-Decoder Paradigm

How to integrate  $(h_1, \dots, h_T)$  in the *decoder* ?

→ It depends what architecture is chosen for the *encoder* and the *decoder*

# The Encoder-Decoder Paradigm

How to integrate  $(h_1, \dots, h_T)$  in the *decoder* ?

- It depends what architecture is chosen for the *encoder* and the *decoder*
- **RNN encoder-decoder** (possibly with an Attention Mechanism)
  - **Transformer Model**

# The Encoder-Decoder with RNNs

**Recall:** Vanilla RNN with  $L'$  and time step sequence of length  $T'$

$$h_{i+1,t+1} = \varphi_i(W_i h_{i,t} + U_i h_{i+1,t} + b_i), \forall i \in [|1, L' - 1|]$$

with  $h_{1,t} = Emb(y_t)$  and  $p_{t+1}^{\hat{}} = h_{L',t+1} \quad \forall t \in [|1, T' - 1|]$

with  $\varphi_{L'} = softmax$

# The Encoder-Decoder with RNNs

Given  $(x_1, \dots, x_T)^T$  and  $(y_1, \dots, y_t) \in V'^{T'}$ , we predict  $\hat{p}_{t+1}$ , distribution over  $V'$  with:

## Decoder

$$h_{dec,i+1,t+1} = \varphi_i(W'_i h_{dec,i,t} + U'_i h_{dec,i+1,t} + b'_i + V_i h_{enc,L+1,T+1}), \forall i \in [1, L']$$

$$\text{with } h_{dec,1,t} = Emb(y_t) \text{ and } p_{t+1} = h_{dec,L'+1,t+1} \quad \forall t \in [1, T']$$

$$\text{with } \varphi_{L'} = softmax$$

# The Encoder-Decoder with RNNs

Given  $(x_1, \dots, x_T)^T$  and  $(y_1, \dots, y_t) \in V'^{T'}$ , we predict  $\hat{p}_{t+1}$ , distribution over  $V'$  with:

## Decoder

$$h_{dec,i+1,t+1} = \varphi_i(W'_i h_{dec,i,t} + U'_i h_{dec,i+1,t} + b'_i + V_i h_{enc,L+1,T+1}), \forall i \in [1, L']$$

with  $h_{dec,1,t} = Emb(y_t)$  and  $p_{t+1} = h_{dec,L'+1,t+1} \quad \forall t \in [1, T']$

with  $\varphi_{L'} = softmax$

- Decoder of *with  **$L'$  layer***, with parameters  **$W'_i, U'_i, b', V_i$  for all  $i$**
- It decodes sequentially the target sequence



# The Encoder-Decoder with RNNs

Given  $(x_1, \dots, x_T)^T$  and  $(y_1, \dots, y_t) \in V'^{T'}$ , we predict  $\hat{p}_{t+1}$ , distribution over  $V'$  with:

## Decoder

$$h_{dec,i+1,t+1} = \varphi_i(W'_i h_{dec,i,t} + U'_i h_{dec,i+1,t} + b'_i + V_i h_{enc,L+1,T+1}), \forall i \in [1, L']$$

$$\text{with } h_{dec,1,t} = Emb(y_t) \text{ and } p_{t+1} = h_{dec,L'+1,t+1} \quad \forall t \in [1, T']$$

$$\text{with } \varphi_{L'} = softmax$$

- Decoder of *with  $L'$  layer*, with parameters  $W'_i, U'_i, b', V_i$  for all  $i$
- It decodes sequentially the target sequence
- It is conditioned on **the input sequence through the encoding output**

# The Encoder-Decoder with RNNs

Given  $(x_1, \dots, x_T)^T$  and  $(y_1, \dots, y_t) \in V'^{T'}$ , we predict  $\hat{p}_{t+1}$ , distribution over  $V'$  with:

## Decoder

$$h_{dec,i+1,t+1} = \varphi_i(W'_i h_{dec,i,t} + U'_i h_{dec,i+1,t} + b'_i + V_i h_{enc,L+1,T+1}), \forall i \in [1, L']$$

$$\text{with } h_{dec,1,t} = Emb(y_t) \text{ and } p_{t+1} = h_{dec,L'+1,t+1} \quad \forall t \in [1, T']$$

$$\text{with } \varphi_{L'} = softmax$$

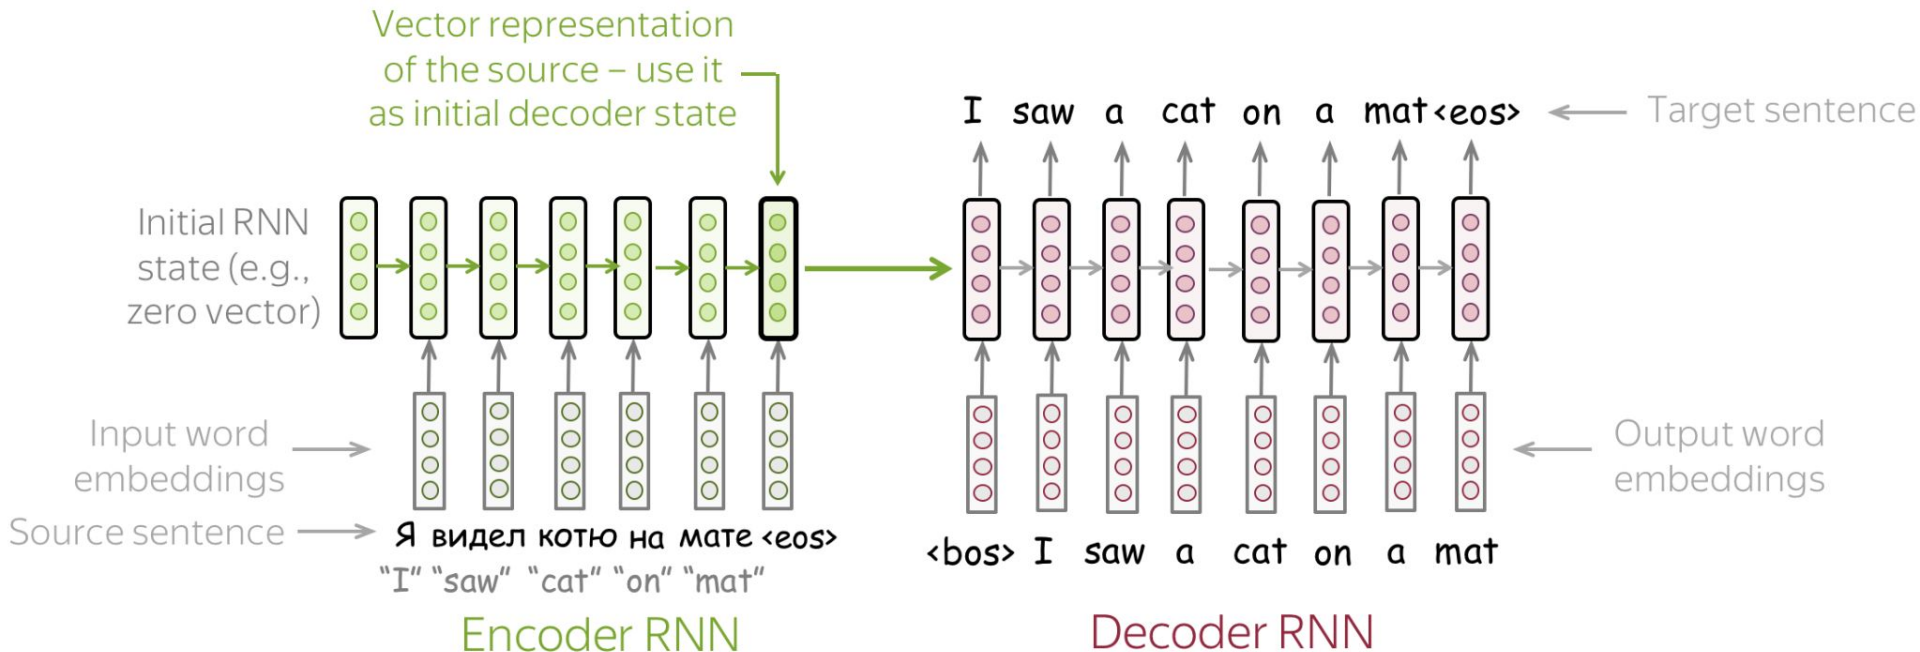
**Encoder: also a RNN that encodes the input sequence in a fixed vector**

$$h_{enc,i+1,t+1} = \varphi_i(W_i h_{enc,i,t} + U_{enc,i} h_{i+1,t} + b_{enc,i})$$

$$\forall i \in [1, L] \quad \forall t \in [1, T] \quad \text{with } h_{enc,1,t} = Emb(x_t)$$

# The Encoder-Decoder with RNNs

## Simple RNN-based Encoder-Decoder Model:



# The Encoder-Decoder Training

## With Backpropagation

1. We feed the model with both the input and output sequence
2. We compute the loss based on the “gold” output sequence
3. We update all the parameters of the model with backpropagation

# The Encoder-Decoder with RNNs: Limits

**Limits:** *At test time in the encoder-decoder that we introduced*  
The input sequence has a fixed representation ( $h_{enc,L+1,T+1}$  is fixed)

# The Encoder-Decoder with RNNs: Limits

**Limits:** *At test time in the encoder-decoder that we introduced*  
The input sequence has a fixed representation ( $h_{enc,L+1,T+1}$  is fixed)

**Example:**

*Je vois un chat sur un matelas*  $\Rightarrow$  *I see a cat on a mat*

# The Encoder-Decoder with RNNs: Limits

**Limits:** *At test time in the encoder-decoder that we introduced*  
The input sequence has a fixed representation ( $h_{enc,L+1,T+1}$  is fixed)

**Example:**

*Je vois un chat sur un matelas*  $\Rightarrow$  *I see a cat on a mat*

**Step 4:**

**Given:** *Je vois un chat sur un matelas*  $\Rightarrow$  *cat*

**Step 7:**

**Given:** *Je vois un chat sur un matelas*  $\Rightarrow$  *mat*

# The Encoder-Decoder with RNNs: Limits

**Limits:** *At test time in the encoder-decoder that we introduced*  
The input sequence has a fixed representation ( $h_{enc,L+1,T+1}$  is fixed)

**Example:**

*Je vois un chat sur un matelas*  $\Rightarrow$  *I see a cat on a mat*

**Step 4:**

**Given:** *Je vois un chat sur un matelas*  $\Rightarrow$  *cat*

**Step 7:**

**Given:** *Je vois un chat sur un matelas*  $\Rightarrow$  *mat*

$\rightarrow$  We need “decoding-dependent” representation of the input sequence 56



# How to build more flexible encoder-decoder?

## Solution 1:

- Integrate an Attention Mechanism in a RNN-based encoder-decoder

## Solution 2:

- Use an Encoder-Decoder Transformer

In lab you will only make a classic seq2seq!

## To Be Seen in Lab

- Classic Unsupervised learning to explore data (K-Means)
- PyTorch tutorial
- Unsupervised learning with PyTorch (Autoencoder)